

Quillian

Private by design. Powerful by default.

Quillian is a privacy-native document intelligence system that lets organizations ask natural language questions of their document corpus and receive cited, sourced answers, all without sending a single document to the cloud. Access PDF-locked data via inquiry, not only search.

Unlike cloud-based document AI tools, Quillian runs entirely on your infrastructure. That means you can work with the data you have including protected health information (PHI), personally identifiable data (PII), confidential research, privileged legal documents, and proprietary records, without compromising your data governance obligations.

W H Y P D F

Unlock the Knowledge Already in Your Documents

Most institutional knowledge lives in PDFs, research papers, clinical protocols, policy documents, legal filings, grant reports, meeting minutes, and internal records accumulated over years or decades. These documents are searchable by keyword but not by meaning.

Quillian changes that. Ask "What were the primary risk factors identified across our clinical trial documentation?" and get a cited answer drawn from the relevant passages, without manually reading a single document. Ask "Which policy documents address data retention for PHI?" and receive a sourced response with page-level references.

This is not keyword search. Quillian understands the semantic meaning of your question and matches it against the meaning of your documents, surfacing answers that keyword search would miss entirely.

Because Quillian ingests PDFs directly and indexes them locally, there is no conversion step, no upload, and no third-party processing. Your documents go from folder to queryable corpus in minutes.

T H E P R O B L E M

Cloud RAG Creates Unacceptable Risk for Sensitive Organizations

Retrieval-Augmented Generation (RAG) has become the dominant architecture for document intelligence. However, most RAG implementations are cloud dependent. This means documents are uploaded to external servers, queries are processed by third-party APIs, and data crosses trust boundaries with every interaction.

For organizations operating under HIPAA, IRB protocols, attorney-client privilege, institutional data governance policies, or basic competitive confidentiality requirements, this is not a viable path. Even when a contract specifies that data are protected, the data transfer process itself introduces multiple points of potential breach.

T H E S O L U T I O N

Local-First Document Intelligence, On Your Terms

Quillian gives organizations full control over their data pipeline:

- **Fully Local.** All processing runs on your hardware. No external API calls. No data leaves your network. Compatible with air-gapped environments.
- **Hybrid.** Local retrieval and indexing with optional cloud LLM integration for answer synthesis, on your terms, with your API keys. Shared data are limited, and you choose the sharing scope.
- **Public Sources.** Connect to public institutional repositories (DSpace, PubMed) for large-scale document ingestion without privacy concerns.

H O W I T W O R K S

The Quillian Pipeline

- **Ingest.** PDFs are ingested from a local folder or institutional repository. Text is extracted page by page using pdfplumber, no conversion, no upload, no intermediate processing.
- **Chunk.** Documents are split into overlapping text chunks, preserving context across chunk boundaries.

- **Embed.** Chunks are encoded into dense vector representations using sentence-transformers (all-MiniLM-L6-v2), running locally.
- **Index.** Vectors are stored in a FAISS index on disk, fast, local, and persistent across sessions.
- **Retrieve.** Queries are embedded and matched against the index using cosine similarity. Results are reranked using Maximal Marginal Relevance (MMR) to maximize diversity and eliminate duplicate sources.
- **Synthesize.** The most relevant passage from each unique document is extracted and returned with linked citations, page numbers, and relevance indicators.

USE CASES

Built for High-Stakes Data Environments

- **Healthcare & Clinical Research.** Query patient records, clinical recommendations, IRB documentation, and research literature without PHI or PII leaving your network.
- **Public Health Agencies.** Analyze surveillance reports, policy documents, and epidemiological literature across large corpora and over years throughout your PDF legacy.
- **Legal & Compliance.** Search privileged documents, contracts, regulatory filings, and case materials with full confidentiality.
- **Academic & Institutional Research.** Index and query large publication corpora, grant documentation, and institutional records.
- **Insurance & Actuarial.** Analyze policy documents, claims records, and regulatory filings in a privacy-compliant environment.

TECHNICAL STACK

Architecture Overview

Component	Technology
API Layer	FastAPI + Uvicorn
Embeddings	sentence-transformers (all-MiniLM-L6-v2)
Vector Index	FAISS (Facebook AI Similarity Search)
PDF Parsing	pdfplumber
Reranking	Maximal Marginal Relevance (MMR)
Connectors	Local PDF folders, DSpace 7 repositories
External LLM	None required, fully local by default
Deployment	Python 3.12+, Windows / Linux / macOS

Early Access

Quillian is currently available to a limited number of early access partners. We are particularly interested in working with healthcare, public health, legal, and research institutions where privacy-preserving document intelligence addresses a real operational need.

hello@n61.io · n61.io

© 2026 n61. All rights reserved. Anchorage, Alaska.