SEEK2000

Answers grounded in evidence, not AI memory. Private by design. Evidence by default.

SEEK2000 is a local-first evidence synthesis engine that answers questions using authorized evidence sources rather than relying on AI memory. It retrieves relevant evidence, preserves provenance, synthesizes supported answers, and makes uncertainty explicit.

Unlike cloud-based document AI tools, SEEK2000 can run entirely on your infrastructure. That means organizations can work with protected health information (PHI), personally identifiable information (PII), confidential research, privileged legal documents, and proprietary records without surrendering control of their data.

YOUR EVIDENCE, UNLOCKED

Ask Questions of the Knowledge You Already Have

Organizations accumulate enormous amounts of knowledge in PDFs, research papers, clinical protocols, policy documents, legal filings, grant reports, meeting minutes, and internal records. Much of that knowledge can be found by keyword, but not interrogated by meaning or synthesized across sources.

SEEK2000 changes that. Ask, “What were the primary risk factors identified across our clinical trial documentation?” and receive a supported answer with citations to the relevant evidence. Ask, “Which policy documents address data retention for PHI?” and receive a sourced response with page-level references.

SEEK2000 uses semantic retrieval to identify evidence based on the meaning of a question rather than exact keyword matches. Local documents can be indexed directly without uploading them to a third-party service, turning existing collections into queryable evidence corpora while preserving organizational control of the source material.

THE PROBLEM

Cloud RAG Creates Risk for Sensitive Organizations

Retrieval-Augmented Generation (RAG) has become a common architecture for document intelligence. Many implementations, however, depend on cloud infrastructure: documents are uploaded to external servers, queries are processed through third-party APIs, and organizational data cross trust boundaries during use.

For organizations operating under HIPAA, IRB protocols, attorney-client privilege, institutional data-governance policies, or competitive confidentiality requirements, those dependencies can create unnecessary risk and limit which information can safely be used.

THE SOLUTION

Evidence-Bounded AI, On Your Terms

SEEK2000 gives organizations control over both their evidence and the systems used to analyze it:

Fully Local. Processing can run exclusively on organizational hardware, with no external API calls and no source documents leaving the network. Compatible with air-gapped environments.

Hybrid. Retrieval and indexing remain local while organizations may optionally use external LLMs for synthesis using their own API credentials and policies. Organizations determine what information, if any, crosses the boundary.

Multiple Evidence Sources. SEEK2000 is designed to work across local files, public evidence sources, and external systems through native connectors and standards-based interfaces such as the Model Context Protocol (MCP). This allows new evidence sources to be incorporated without making any single vendor or API the center of the architecture.

Evidence Before Inference. Answers are grounded in retrieved evidence rather than AI memory. SEEK2000 constrains unsupported generation by preserving the evidentiary basis of responses and making uncertainty or insufficient evidence explicit.

HOW IT WORKS

The SEEK2000 Pipeline

Ingest. Evidence is ingested from authorized sources. The current implementation supports local PDF collections, with additional evidence connectors planned.

Chunk. Documents are divided into overlapping text segments to preserve context across boundaries.

Embed. Content is encoded into dense vector representations locally using sentence-transformer models.

Index. Vector representations are stored in a persistent local FAISS index for rapid retrieval.

Retrieve. Questions are semantically matched against indexed evidence. Results are reranked using Maximal Marginal Relevance (MMR) to improve diversity and reduce redundant sources.

Synthesize. Retrieved evidence is assembled into a supported response with linked citations, page-level provenance, and explicit indication of uncertainty or insufficient evidence.

USE CASES

Built for High-Stakes Data Environments

Healthcare & Clinical Research. Query clinical documentation, research literature, protocols, and other authorized evidence while keeping PHI and PII under organizational control.

Public Health Agencies. Analyze surveillance reports, policy documents, epidemiological literature, and institutional records across large historical collections.

Legal & Compliance. Analyze privileged documents, contracts, regulatory filings, policies, and case materials while maintaining control over confidential source material.

Academic & Institutional Research. Query publication corpora, grant documentation, research records, policies, and institutional knowledge.

Insurance & Actuarial. Analyze policy documents, claims records, regulatory materials, and other authorized evidence within controlled environments.

TECHNICAL STACK

Architecture Overview

Component	Technology
API Layer	FastAPI + Uvicorn
Embeddings	sentence-transformers (all-MiniLM-L6-v2)
Vector Index	FAISS (Facebook AI Similarity Search)
PDF Parsing	pdfplumber

Reranking	Maximal Marginal Relevance (MMR)
Connectors	Local PDF collections, PubMed
Planned Connectors	arXiv, websites
External LLM	None required, fully local by default
Interoperability	Model Context Protocol (MCP)
Deployment	Python 3.12+, Windows / Linux / macOS

Early Access

SEEK2000 is currently available to a limited number of early-access partners. We are particularly interested in working with healthcare, public health, legal, and research organizations where evidence synthesis, privacy, provenance, and organizational control address a meaningful operational need.

hello@n61.io · n61.io